

---

# Text Mining

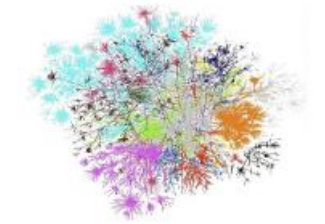
Summer term 2025

Sandipan Sikdar



# Word embeddings

## “Complex Data”



labore et dolore magna aliquyam erat, sed diam voluptua. At vero eos et accusam et justo duo dolores et ea rebum. Stet clita kasd gubergren, no sea takimata sanctus est Lorem ipsum dolor sit amet. Lorem ipsum dolor sit amet, consetetur sadipscing elitr, sed diam nonumy eirmod tempor invidunt ut labore et dolore magna aliquyam erat, sed diam voluptua. At vero eos et accusam et justo duo dolores et ea rebum. Stet clita kasd gubergren.



## “Representation”

1.5
0.2
24
16
2.9

Word2vec, GloVe

## Machine learning algorithm



CNNs, RNNs, ...

Prediction,  
recommendation,  
...

# Text Classification

---

- Supervised learning task
- Given a text, map it to a class
- Example: Sentiment classification

😊 The staff were extremely helpful and friendly. Positioning is great, close to bus/train station and within walking distance to all main sights. Had view over square and room was very clean and comfortable.

😞 Was expensive but perhaps this is simply Venice.

# Text Classification

---

- Train data:  $\{(x^i, y^i)\}_{i=1}^N$

$y^i \in \{0,1\}$



The staff were extremely helpful and friendly. Positioning is great, close to bus/train station and within walking distance to all main sights. Had view over square and room was very clean and comfortable.

$x^i$

- Loss: Binary cross-entropy loss

$$J(\theta) = - \sum_{i=1}^N y^i \log(f_{\theta}(x^i)) + (1 - y^i) \log(1 - f_{\theta}(x^i))$$

Prediction model  
with parameters  $\theta$

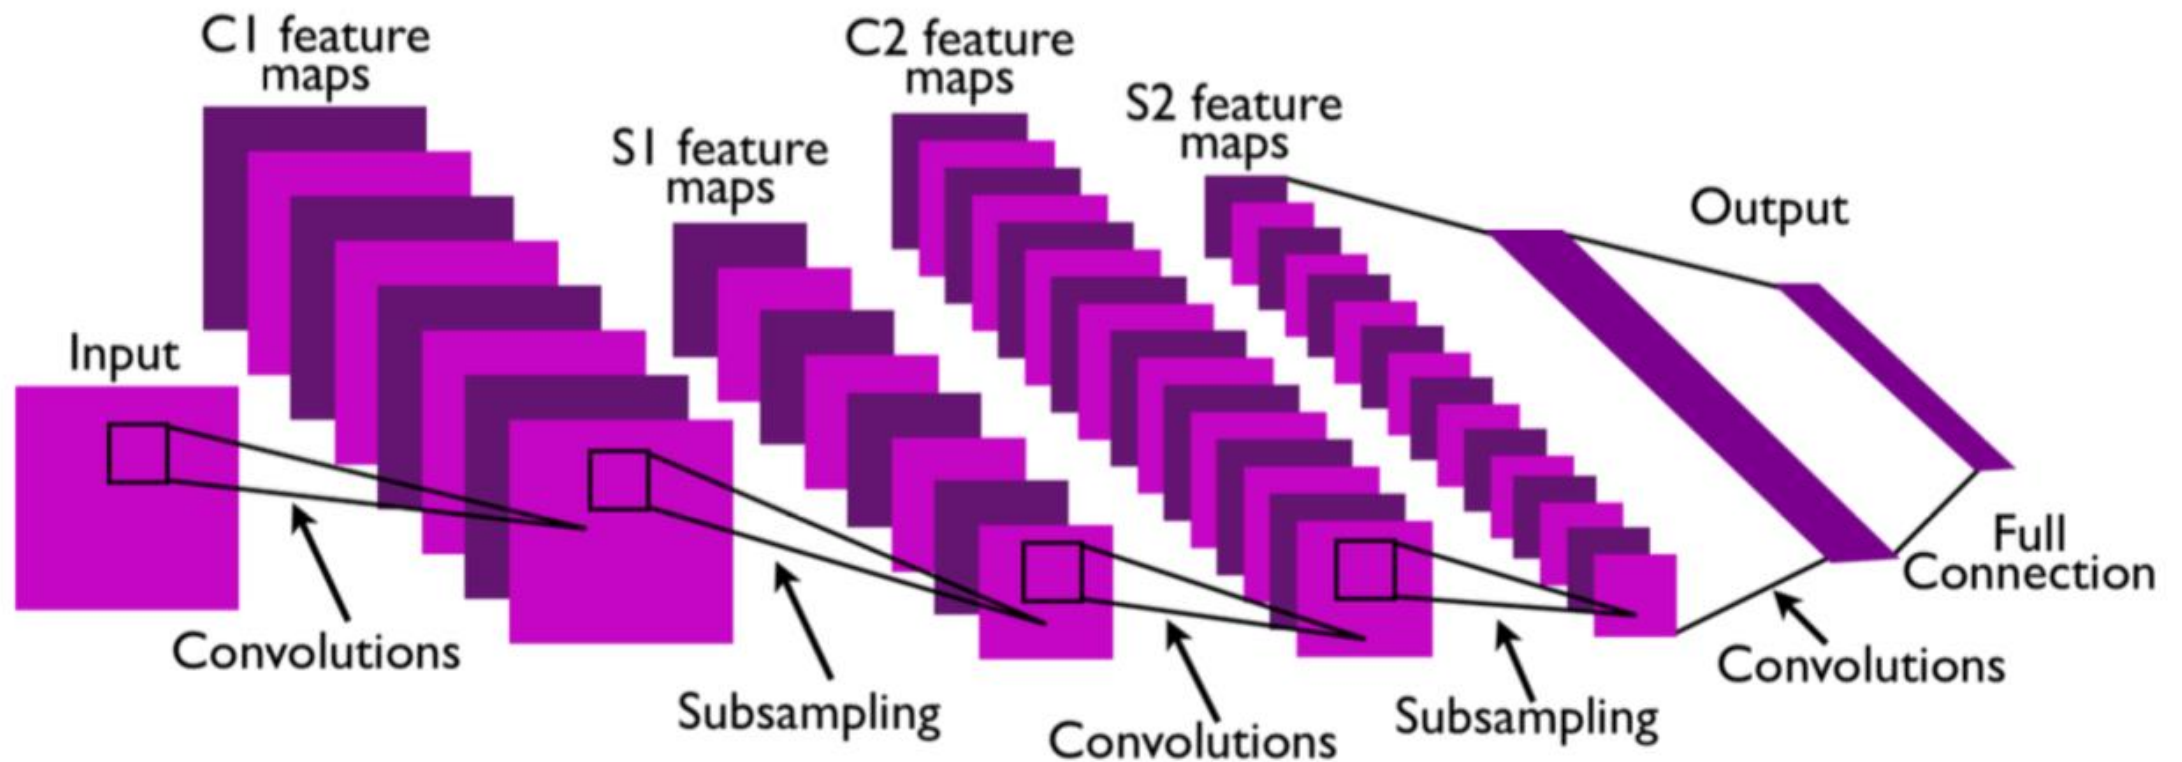
# Convolutional neural networks

---

- ... are a type of partially connected “ feedforward networks” (details later)
- ... originated from computer vision
- ... are often used for image classification
- ... outperform classical image recognition by a large margin
- ... model features over areas of growing size in images
  - Early layers model local neighbourhoods detect edges ,
  - Later layers combine the features to detect larger objects

# Convolutional neural networks

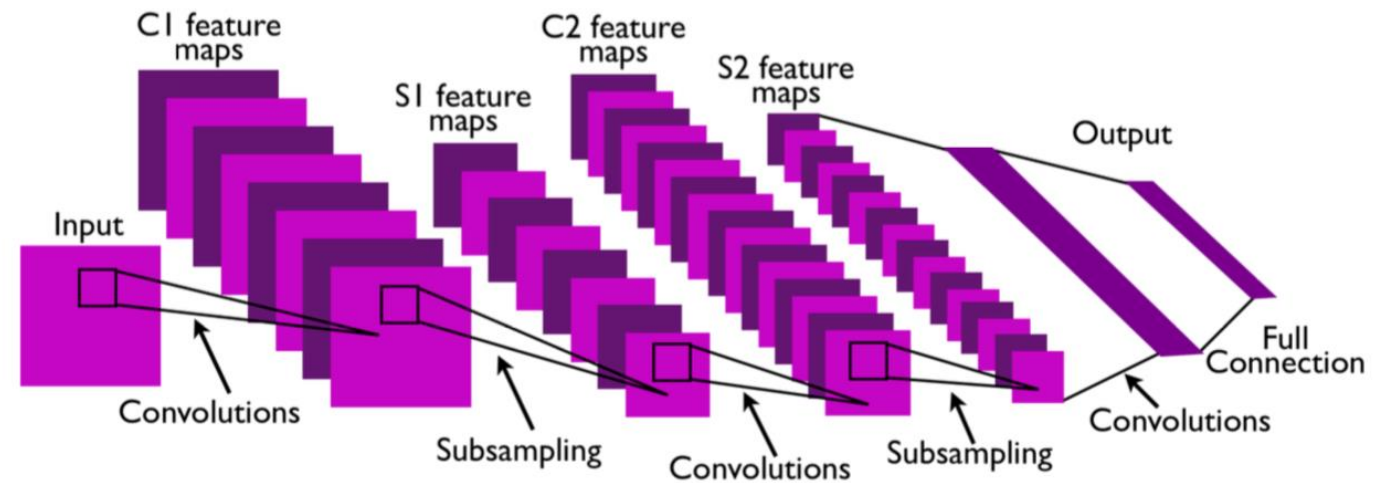
---



<https://wiki.tum.de/display/lfdv/Convolutional+Neural+Networks>

# CNNs in Computer Vision

- Typical structure of CNN
  1. Convolutional layer
  2. Pooling layer
  3. Repeat 1 and 2 multiple times
  4. Fully connected layer
  5. Softmax



# CNN: A bit of history

A bit of history:

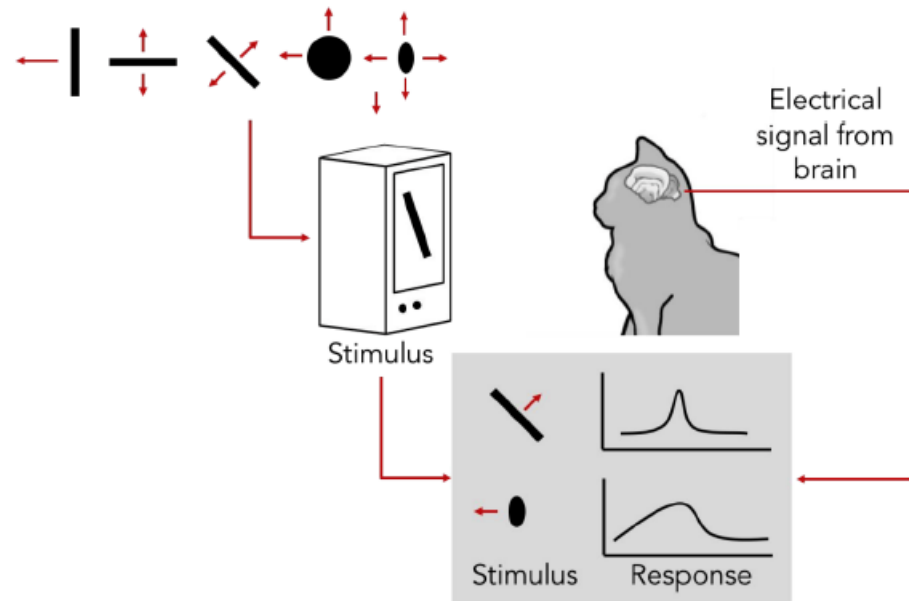
**Hubel & Wiesel,  
1959**

RECEPTIVE FIELDS OF SINGLE  
NEURONES IN  
THE CAT'S STRIATE CORTEX

**1962**

RECEPTIVE FIELDS, BINOCULAR  
INTERACTION  
AND FUNCTIONAL ARCHITECTURE IN  
THE CAT'S VISUAL CORTEX

**1968...**



[Cat image](#) by CNX OpenStax is licensed under CC BY 4.0; changes made



# CNN: A bit of history

## Hierarchical organization

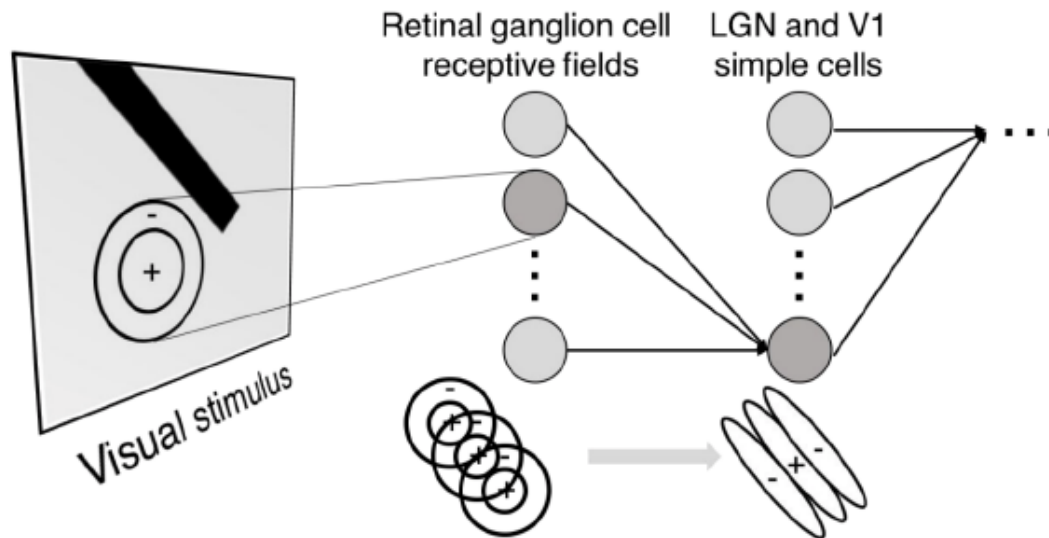


Illustration of hierarchical organization in early visual pathways by Lane McIntosh, copyright CS231n 2017

**Simple cells:**  
Response to light  
orientation

**Complex cells:**  
Response to light  
orientation and movement

**Hypercomplex cells:**  
response to movement  
with an end point



No response



Response  
(end point)

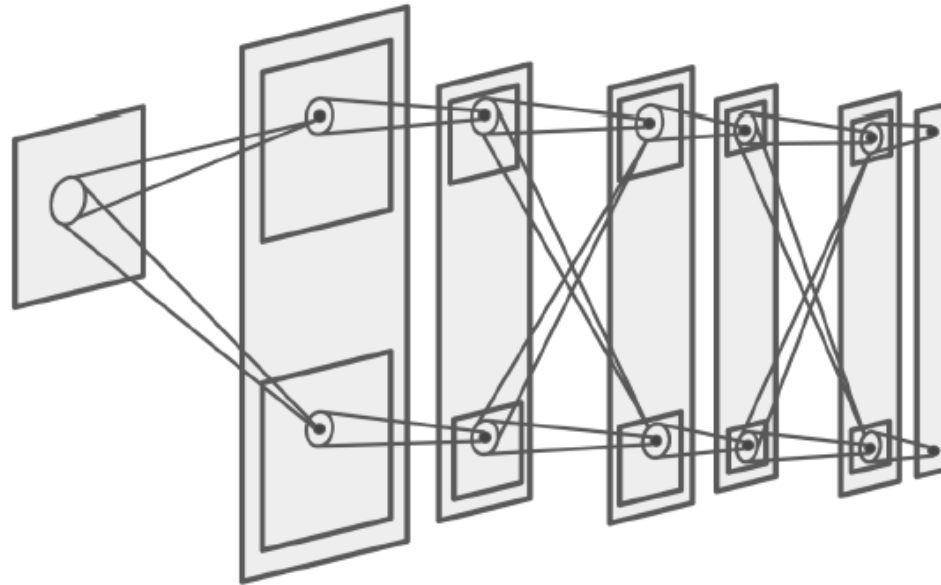
# CNN: A bit of history

---

A bit of history:

**Neocognitron**  
*[Fukushima 1980]*

“sandwich” architecture (SCSCSC...)  
simple cells: modifiable parameters  
complex cells: perform pooling



# CNN: A bit of history

---

"LeNet 1", the first convolutional network that could recognize handwritten digits with good speed and accuracy (1989)

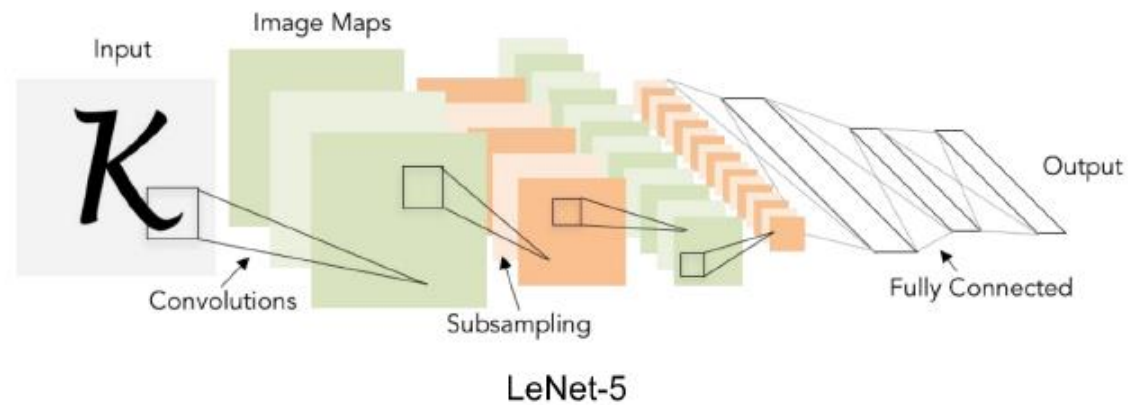


# CNN: A bit of history

---

## Gradient-based learning applied to document recognition

*[LeCun, Bottou, Bengio, Haffner 1998]*



# CNN: A bit of history

---

## ImageNet Classification with Deep Convolutional Neural Networks

*[Krizhevsky, Sutskever, Hinton, 2012]*

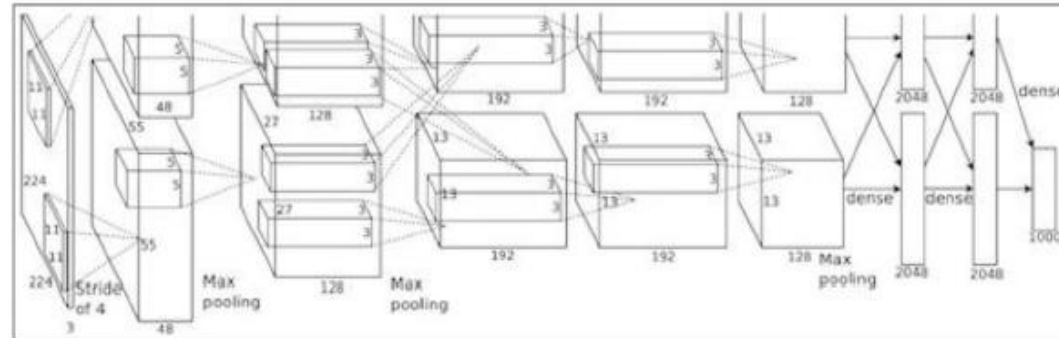


Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

“AlexNet”

# CNNs are everywhere

Classification



Retrieval



Figures copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012.

# Convolutions

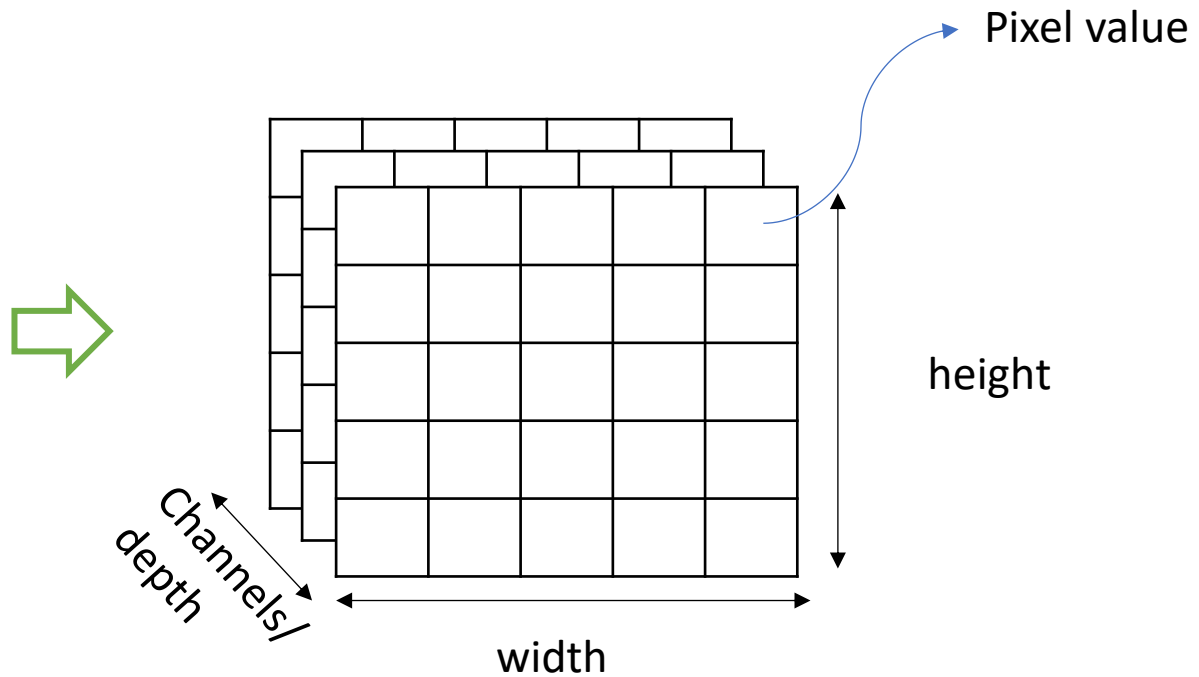
---

- To understand CNNs, we take a look at their basic building blocks  
Convolutions
- Convolution = classic tool in image processing even without machine learning
- Intuition: Slide a function over an image , create a modified image
- Can model operations like edge detection , blurring , sharpening



# Convolutions

- Representing image



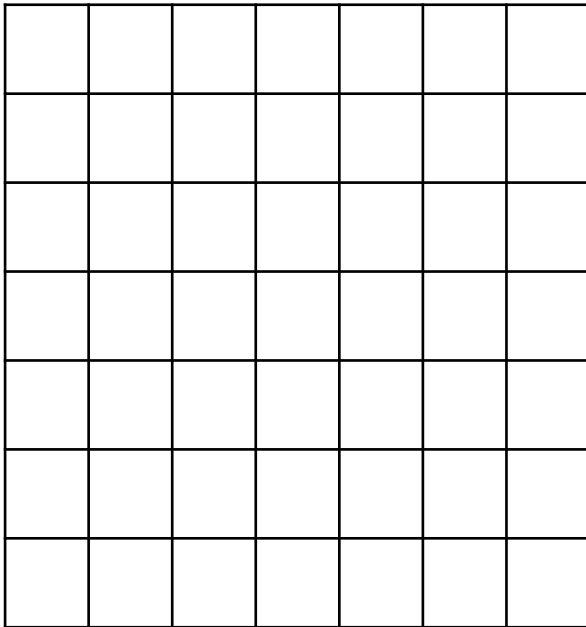


# Convolutions

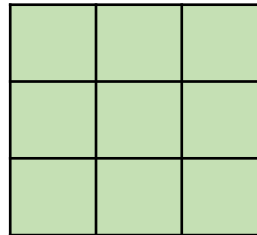
---

We consider a single channel image

Image

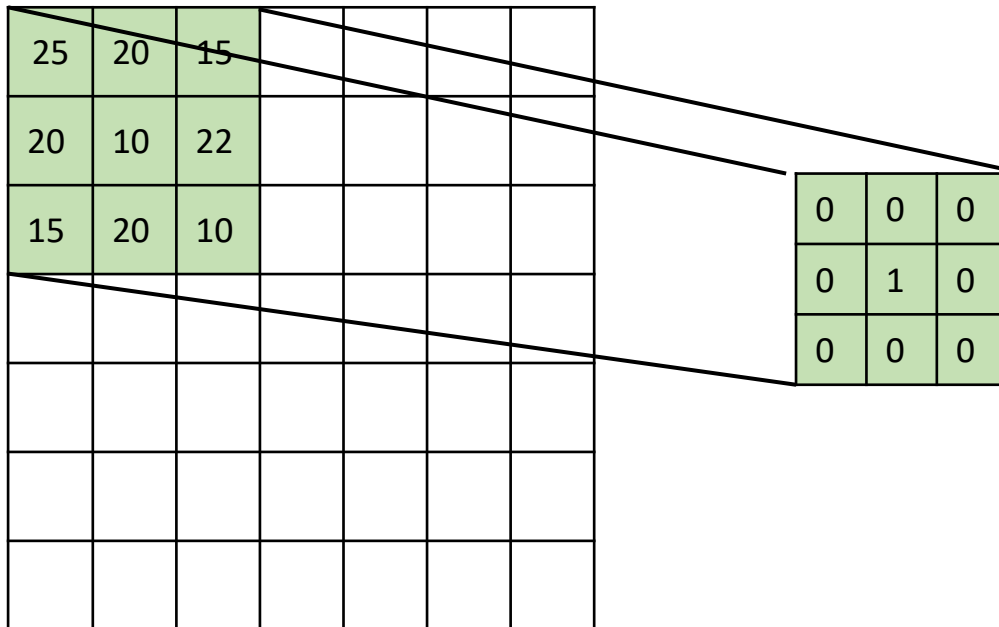


Filter



- Input image (matrix)
- Kernel/filter (matrix)
- Convolve the filter over the image i.e., slide over the image spatially computing dot products

# Convolutions



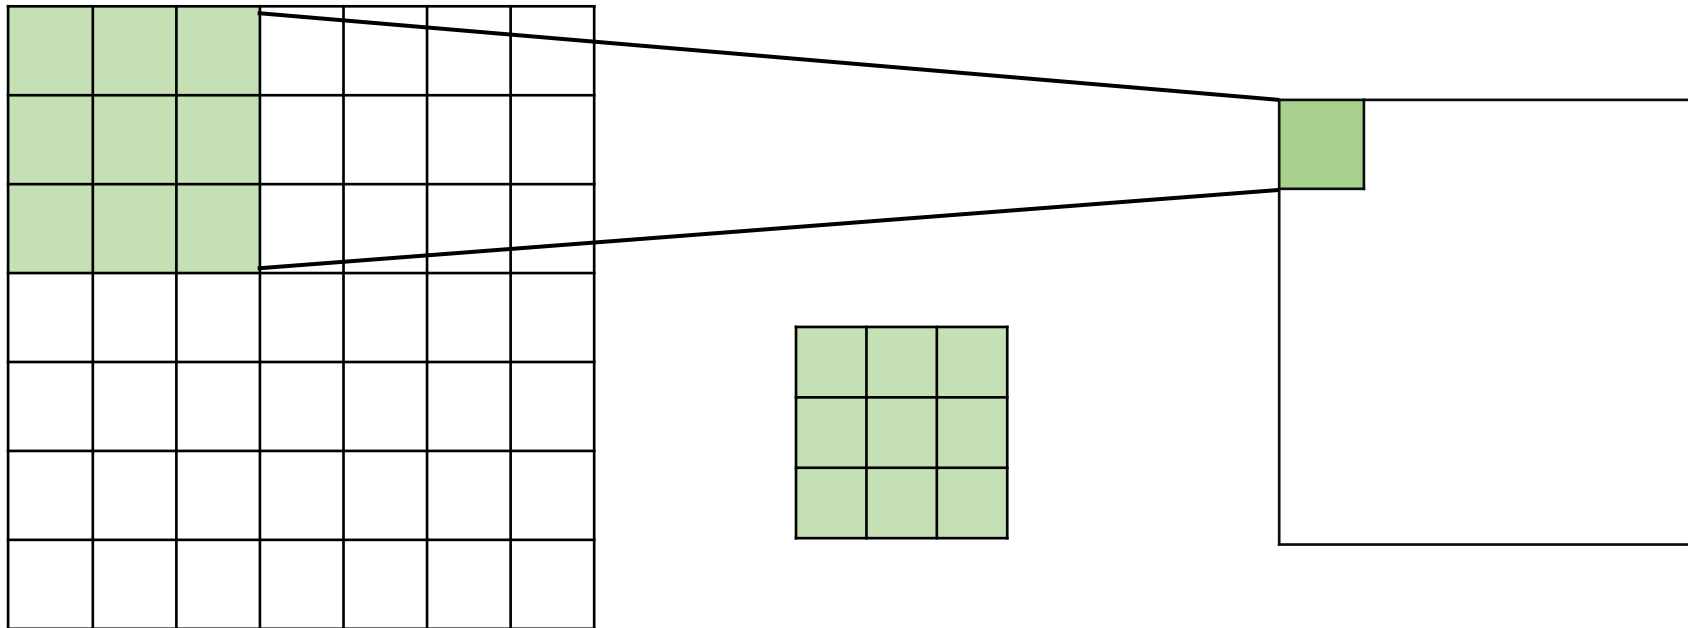
$$out = w^T x + b$$

$$25 * 0 + 20 * 0 + 15 * 0 + 20 * 0 + 10 * 1 + 22 * 0 + 15 * 0 + 20 * 0 + 10 * 0 + 0 = 10$$

(assuming  $b = 0$ )

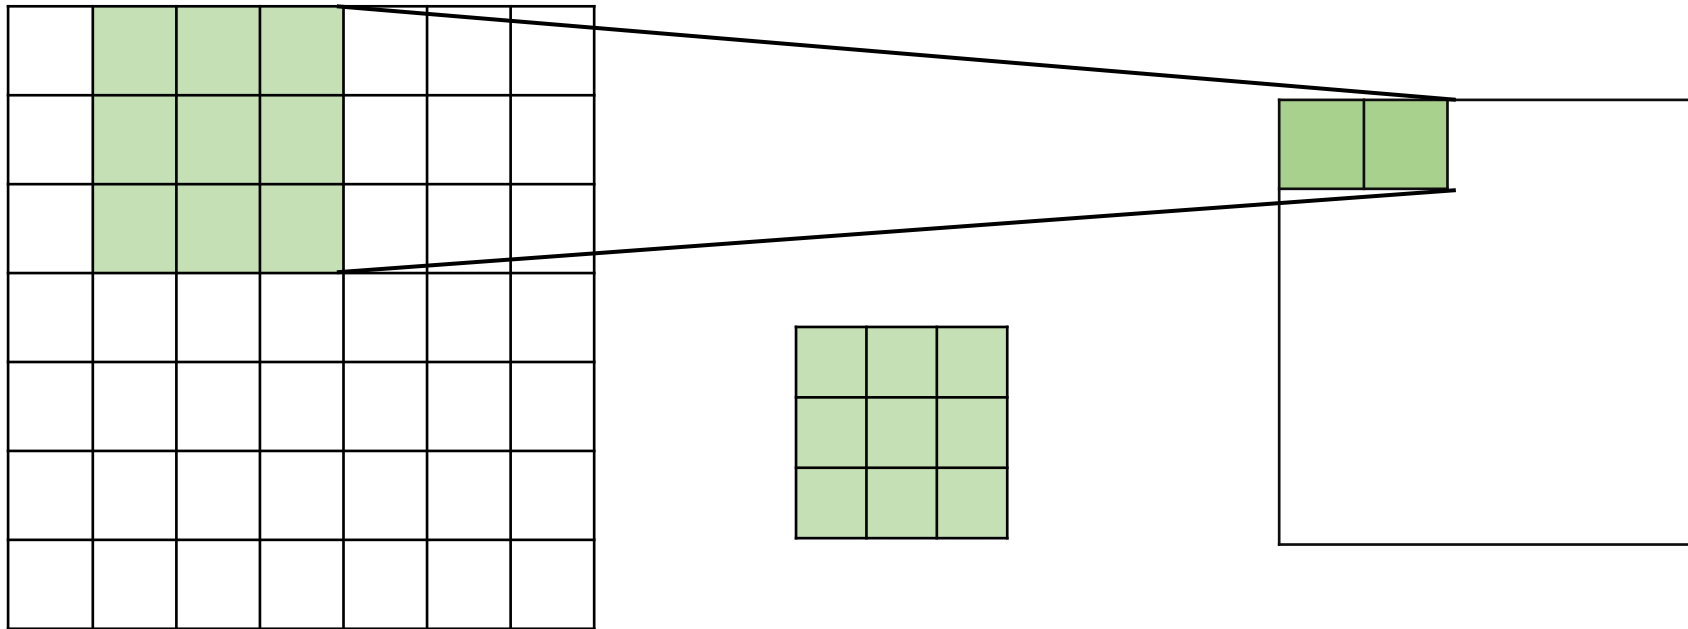
# Convolutions

---



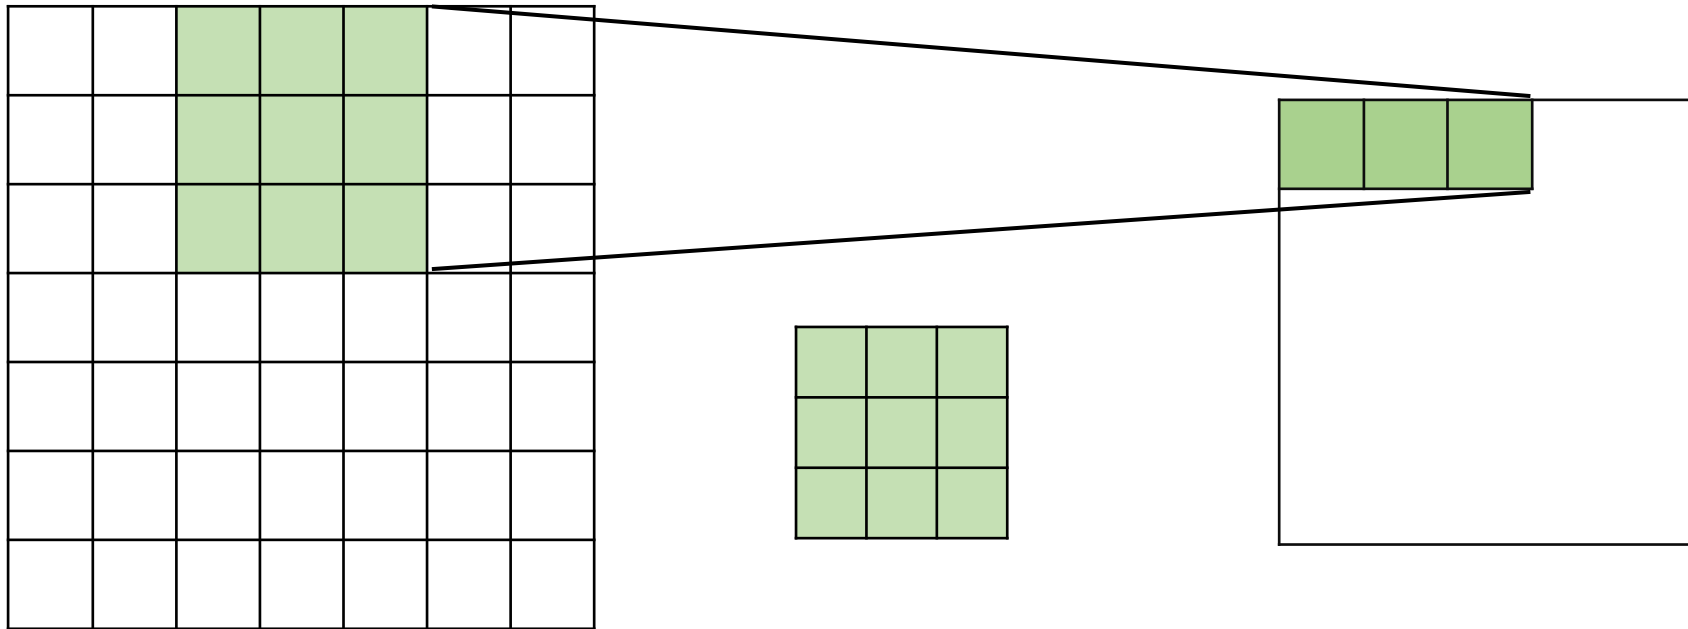
# Convolutions

---



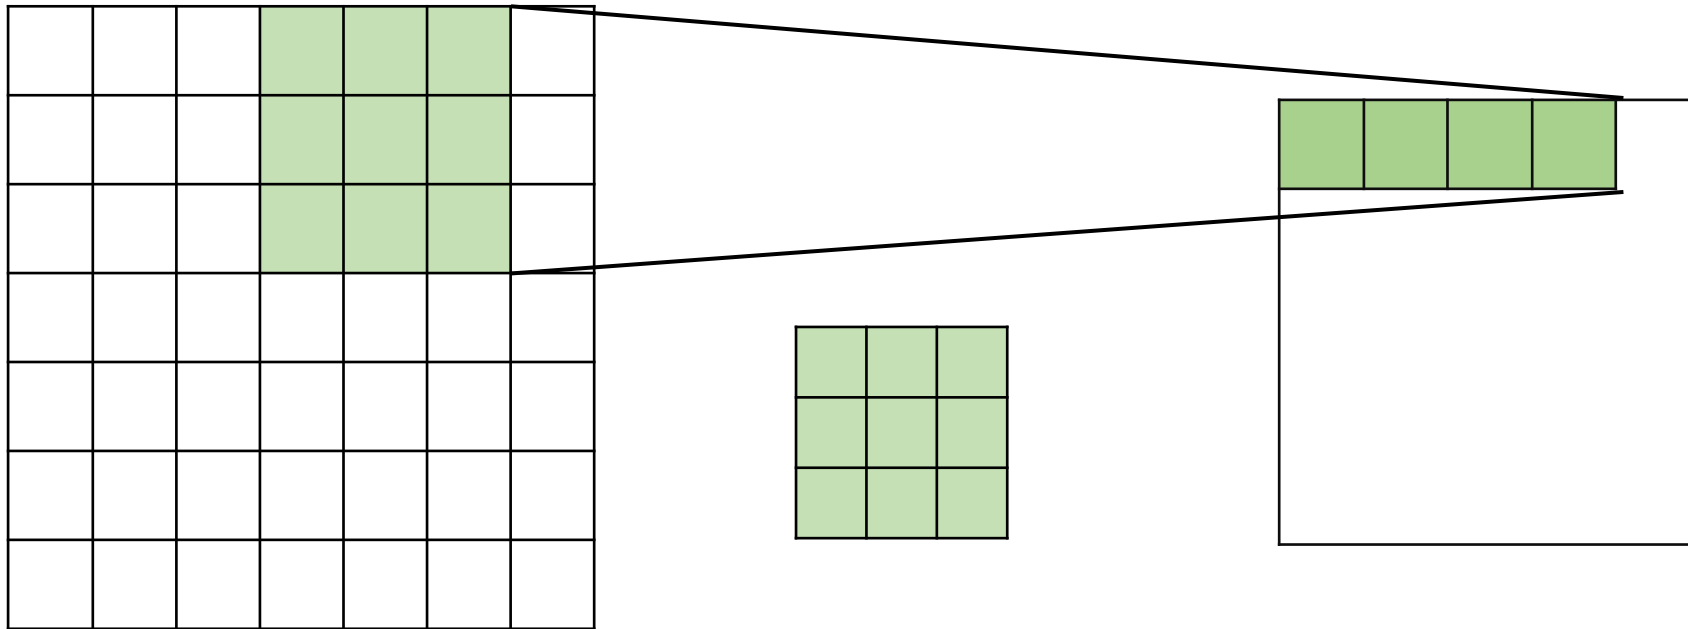
# Convolutions

---



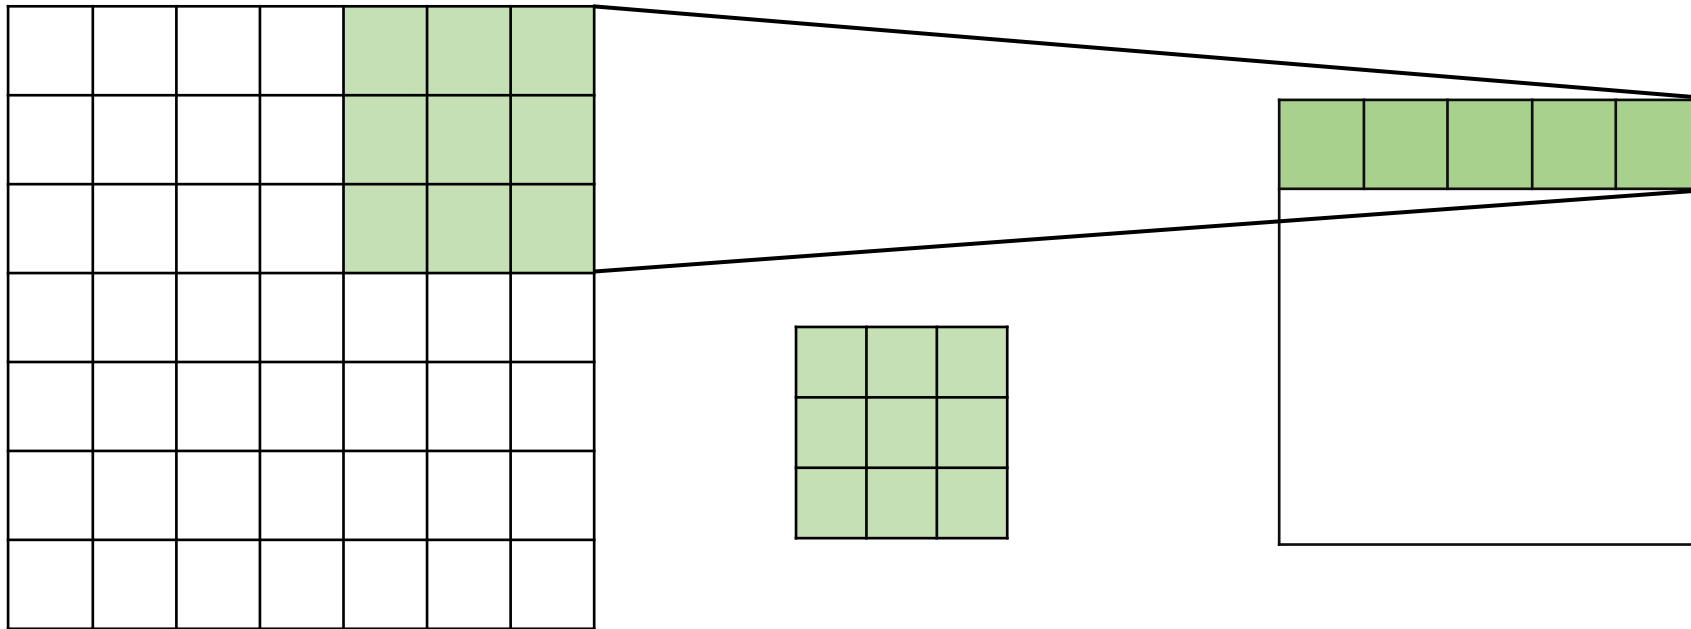
# Convolutions

---



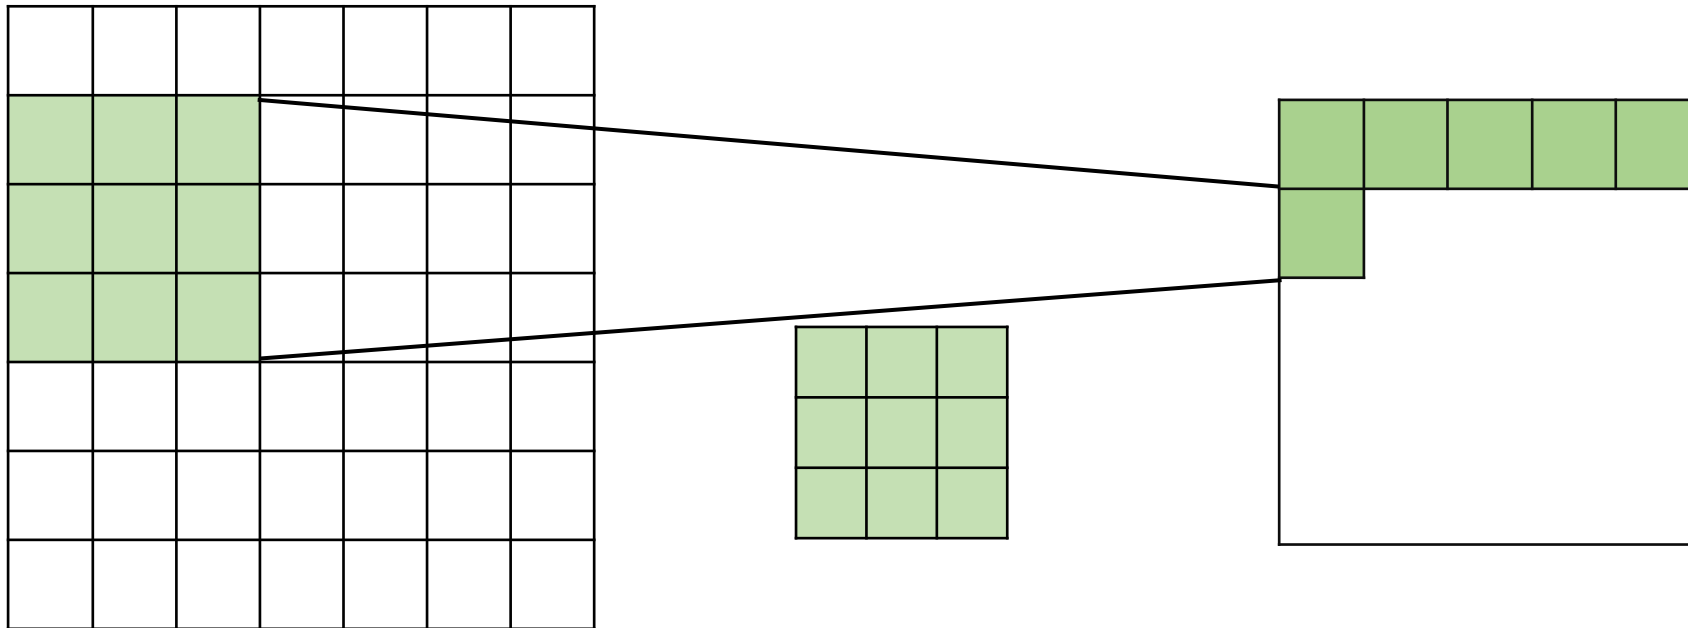
# Convolutions

---



# Convolutions

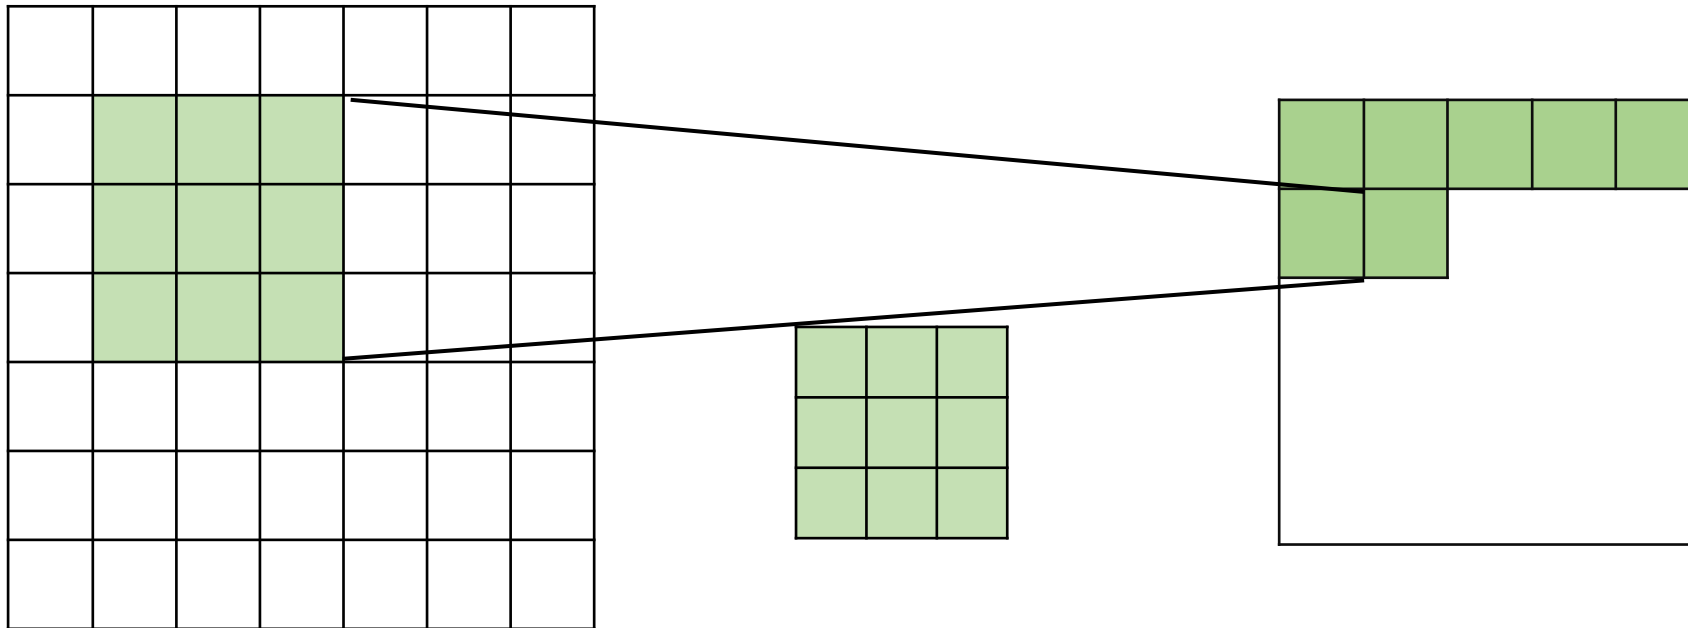
---





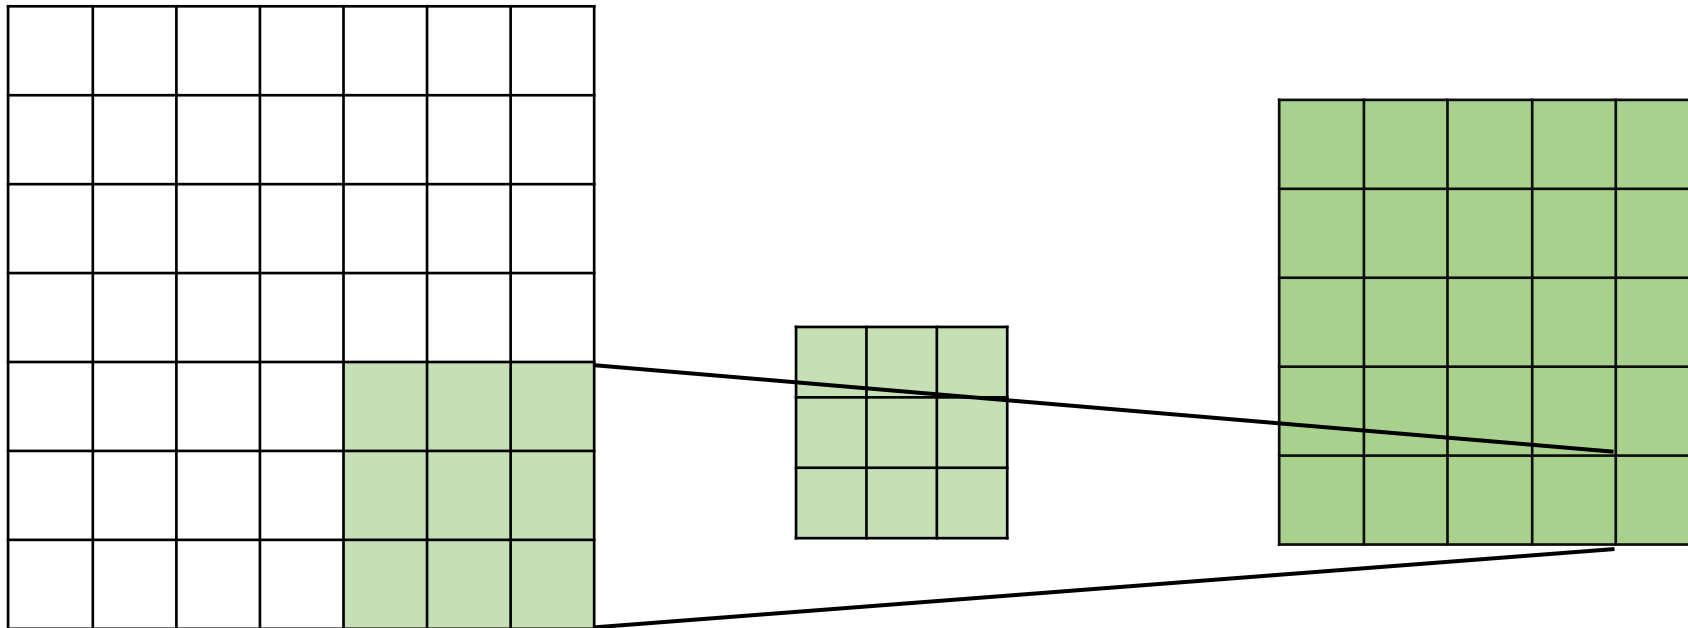
# Convolutions

---



# Convolutions

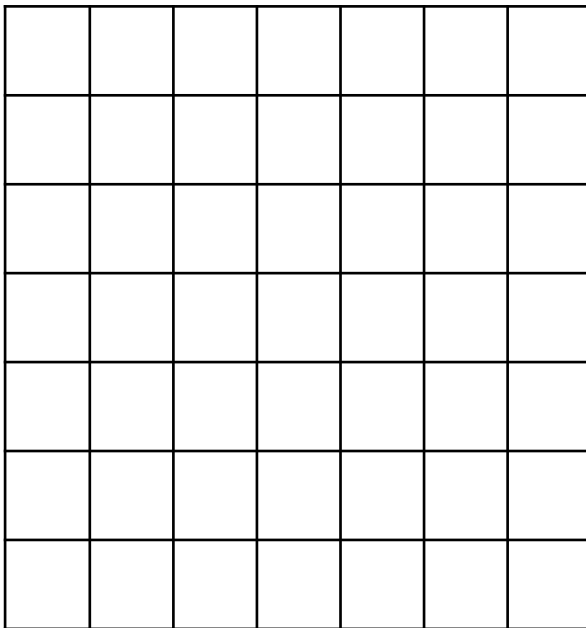
---



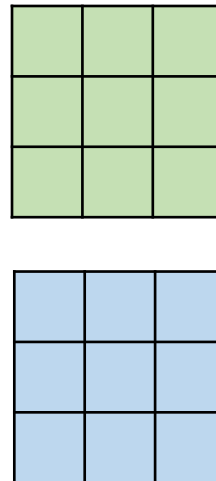
# Convolutions

---

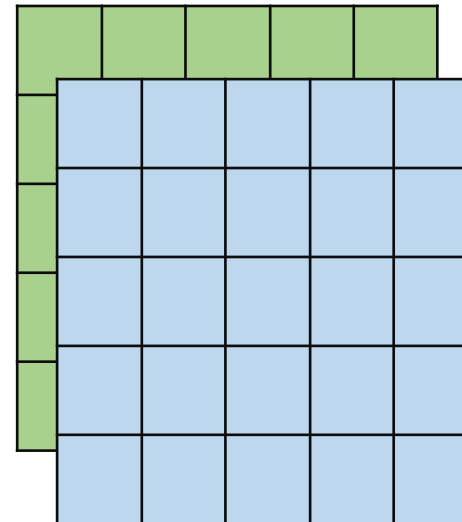
Image



Filters



Feature maps



# What are these feature maps actually?

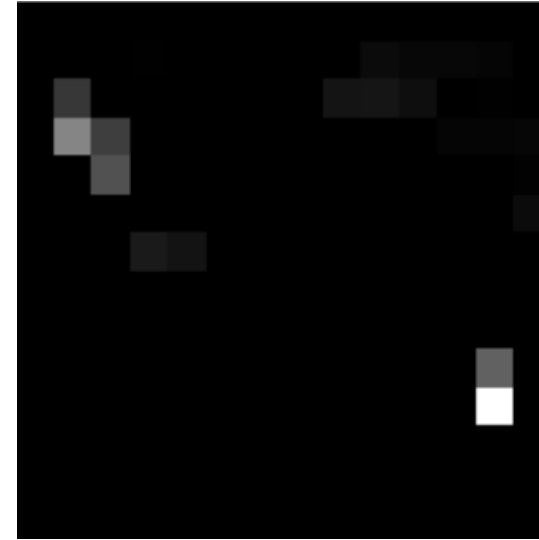
---



Input Image



Feature map obtained  
after first convolution  
operation

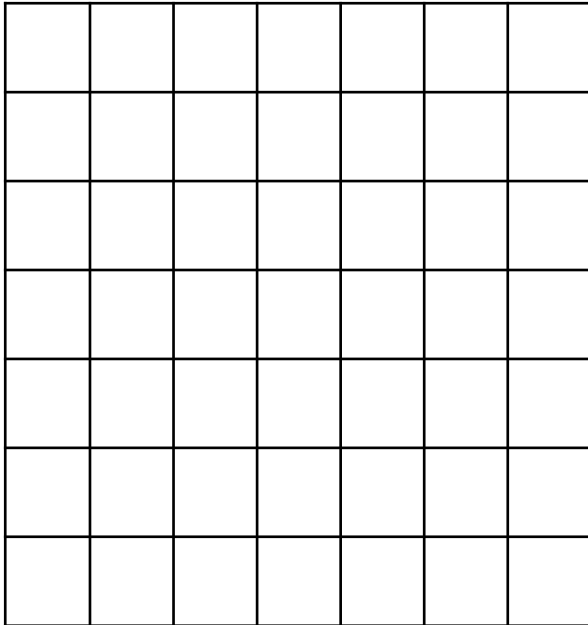


Feature map obtained  
at upper layers

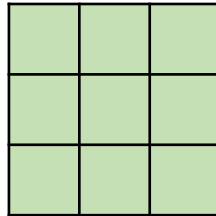
# Convolutions

---

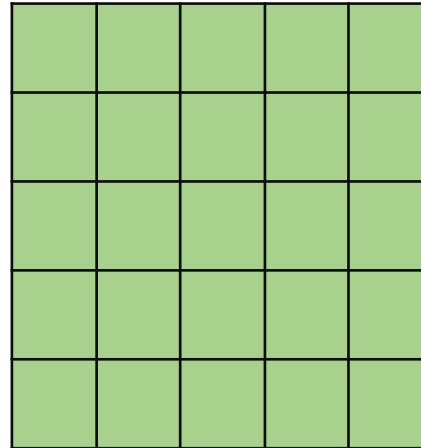
Image



Filter



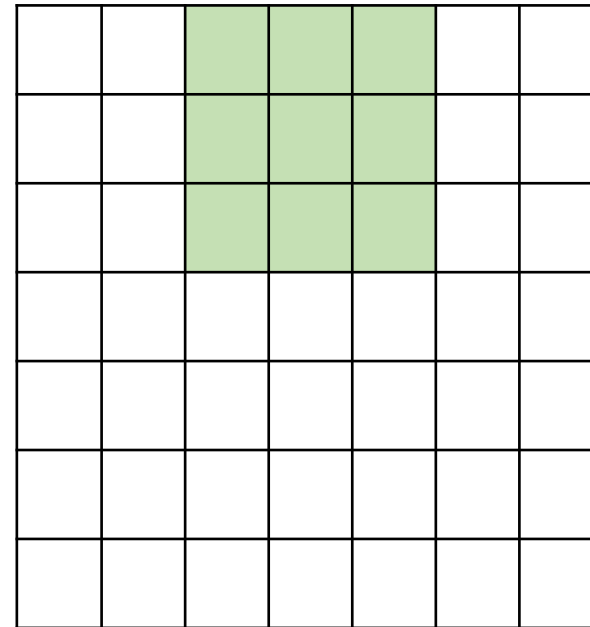
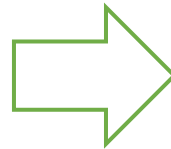
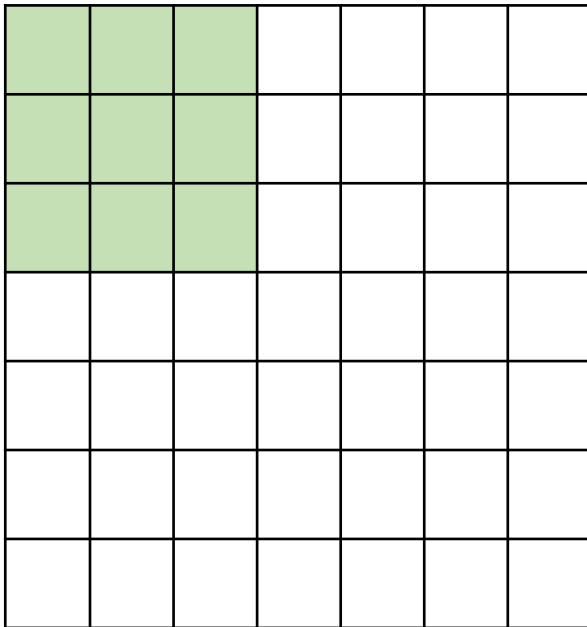
Feature map



- Image dimension: 7x7
- Filter dimension: 3x3
- Stride: 1 (the number of pixels/units) the filter moves towards right or bottom
- Feature map: 5x5

# Convolutions

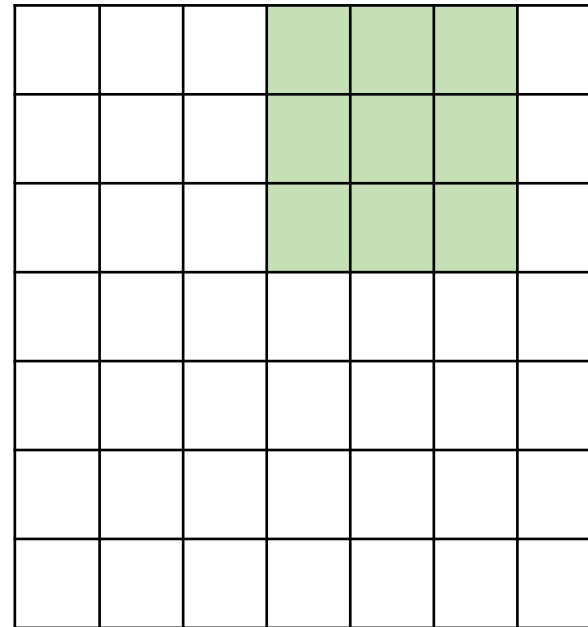
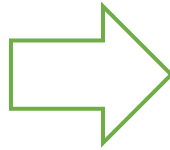
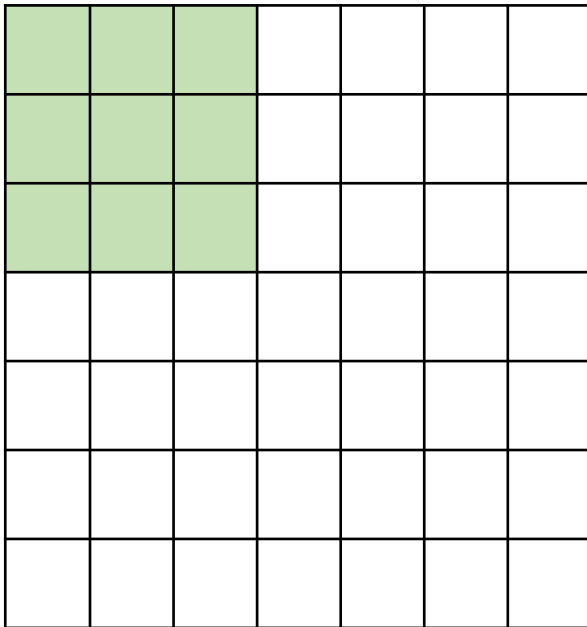
- What is the dimension of the feature map if the stride is 2?



- Feature map: 3x3

# Convolutions

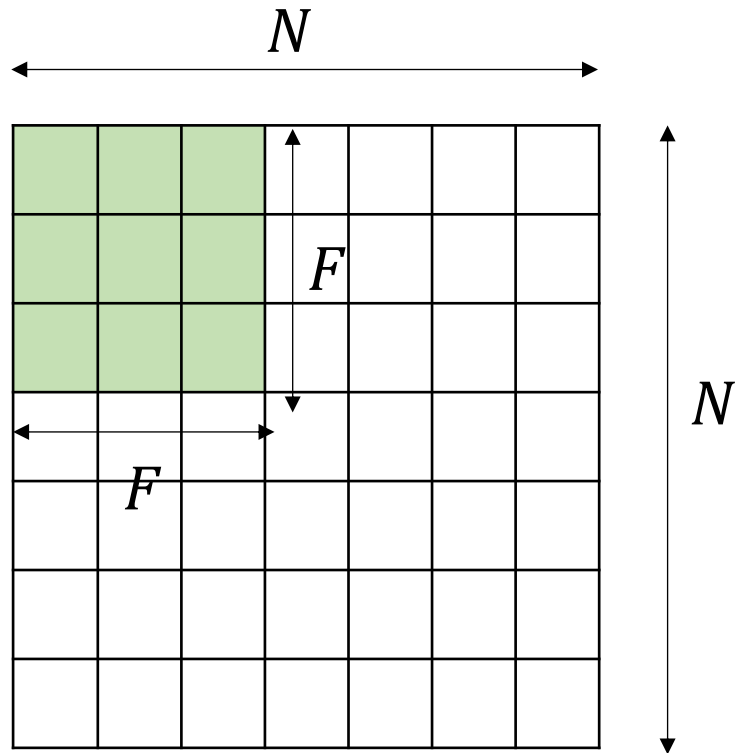
- What happens when the stride is 3?



- Feature map: **Does not fit**

# Convolutions

---



$$\text{Output size} = \frac{N-F}{\text{Stride}} + 1$$

E.g.,

$$N = 7, F = 3, \text{stride} = 1, \text{Output} = 5$$

$$N = 7, F = 3, \text{stride} = 2, \text{Output} = 3$$

$$N = 7, F = 3, \text{stride} = 3, \text{Output} = 2.33$$



# Convolutions

Solution: Zero pad the border

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

$$\text{Output size} = \frac{N-F}{\text{Stride}} + 1$$

E.g.,

Input: 7x7

Filter: 3x3

Stride: 1

Pad with 1 pixel border

Output? 7x7

# Convolutions

Solution: Zero pad the border

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

In general, common to see CONV layers with stride 1, filters of size  $F \times F$ , and zero-padding with  $(F - 1)/2$ . (will preserve size spatially)

E.g.,

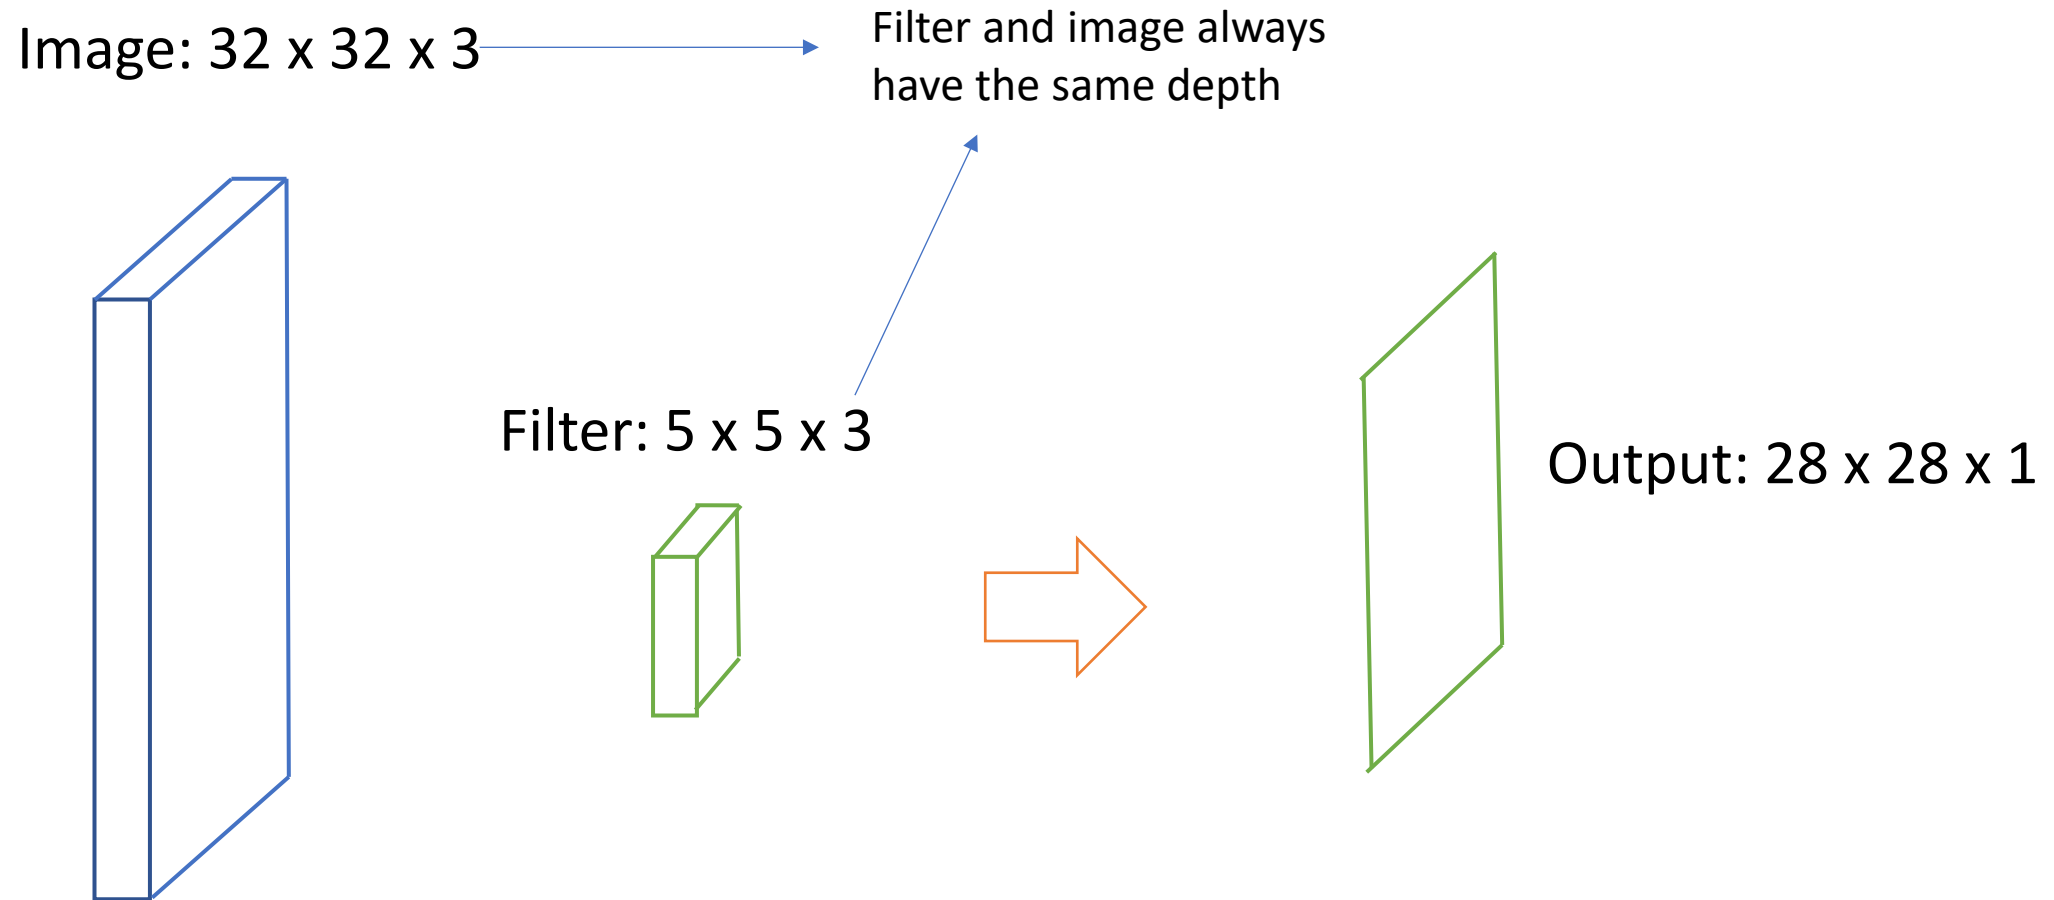
$F = 3 \Rightarrow$  zero pad with 1

$F = 5 \Rightarrow$  zero pad with 2

$F = 7 \Rightarrow$  zero pad with 3

# 3D convolutions

---



# Convolutions

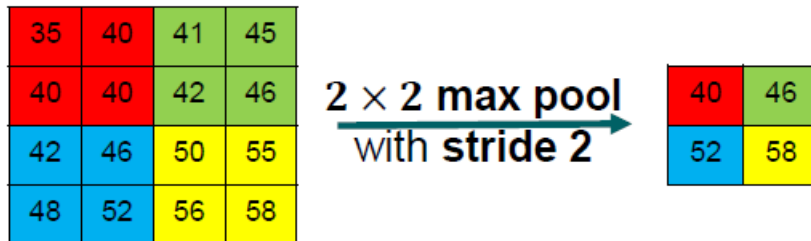
---

- Input:  $W_1 \times H_1 \times D_1$
- Number of filters:  $K$
- Filter dimensions:  $F \times F$
- Stride:  $S$
- Padding:  $P$
- Output
  - $W_2 = \frac{W_1 + 2 * P - F}{S} + 1$
  - $H_2 = \frac{H_1 + 2 * P - F}{S} + 1$
  - $D_2 = K$

Convolution operations are followed by non-linear activations such as Sigmoid or ReLU

# Pooling

- Second component of CNNs: **Pooling Layers**
- Reduce the size of the input image in a predefined manner
- No learned weights ! Purely deterministic operation
- Common types:
  - Max pooling
    - Extract the maximum of  $n \times m$  (pool size) entries in the input
    - Move  $k$  (stride ) entries ahead in the input



Average pooling: Compute the average of inputs instead of max

# Pooling layers

---

- Why use pooling layers?
- Reduce the number of parameters in following layers
- Not all filters activated on all pixels
- Only use pixels in a close neighbourhood that provide the strongest signal

# CNNs vs. Fully connected layers

---

- Massively fewer parameters than fully connected networks
- Example:
  - 300×300 pixel input , map to hidden layer of same size
  - Fully connected :  $(300 \cdot 300) \cdot (300 \cdot 300) = 8,100,000,000$  parameters!
  - Convolutional with 100 filters of size 3×3:  $100 \cdot 3 \cdot 3 = 900$  parameters
- Less RAM needed , less prone to overfitting

---

# CNNs for Text



# CNNs in NLP

---

- So far: Only described use in Computer Vision...
- CNNs became popular in NLP, too:
  - Kalchbrenner et al., 2014: A Convolutional Neural Network for Modelling Sentences
  - Kim, 2014: Convolutional Neural Networks for Sentence Classification
  - Nguyen and Grishman, 2015: Relation Extraction: Perspective from Convolutional Neural Networks
  - ...

# CNNs for Sentence Classification

---

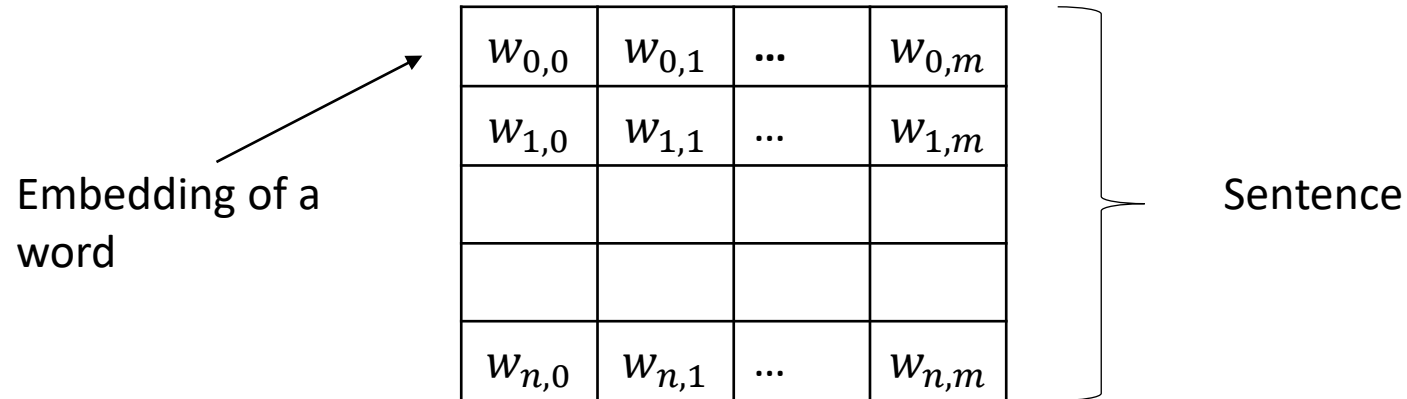
- Kim, 2014 (EMNLP)
- Very simple CNN model...
- ... with very good results!
- Beat previous state-of-the-art in several tasks
  - Sentiment Analysis
  - Subjectivity Detection\*
  - Question Answering

\*didn't beat state-of-the-art, but came very close

# CNNs for Sentence Classification

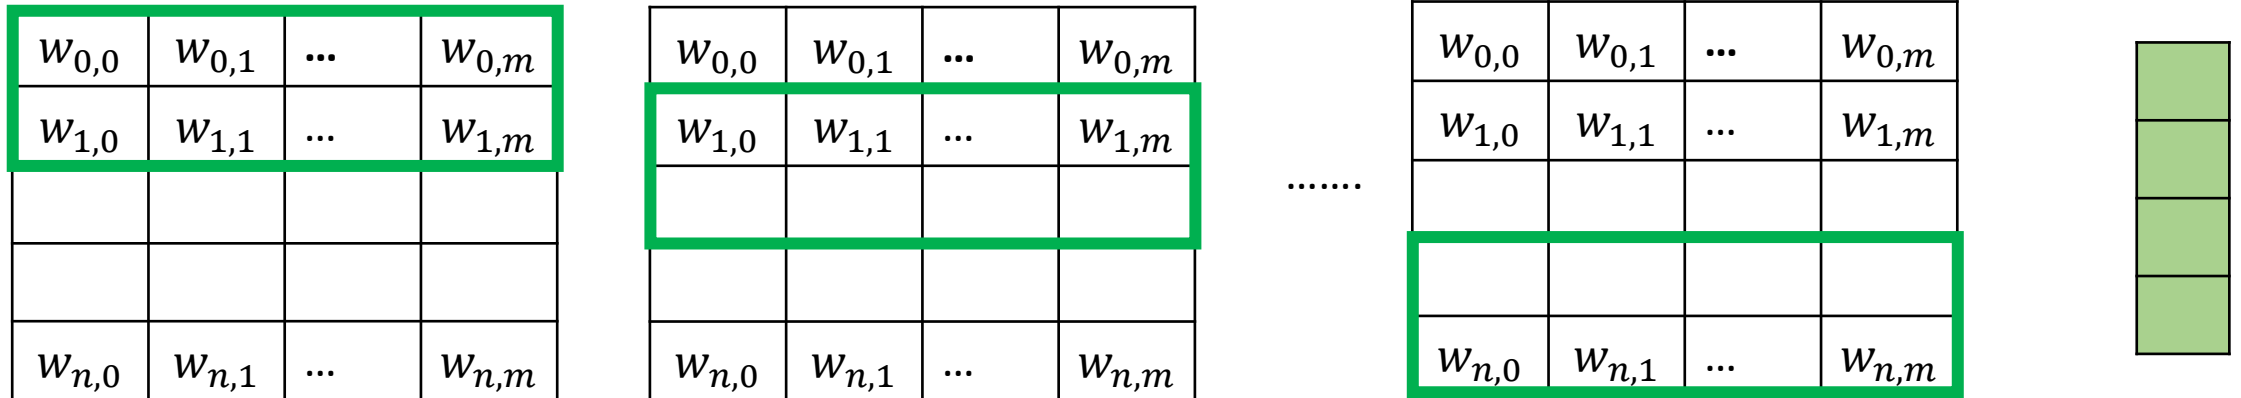
---

- Input representation: Concatenated word embeddings



# CNNs for Sentence Classification

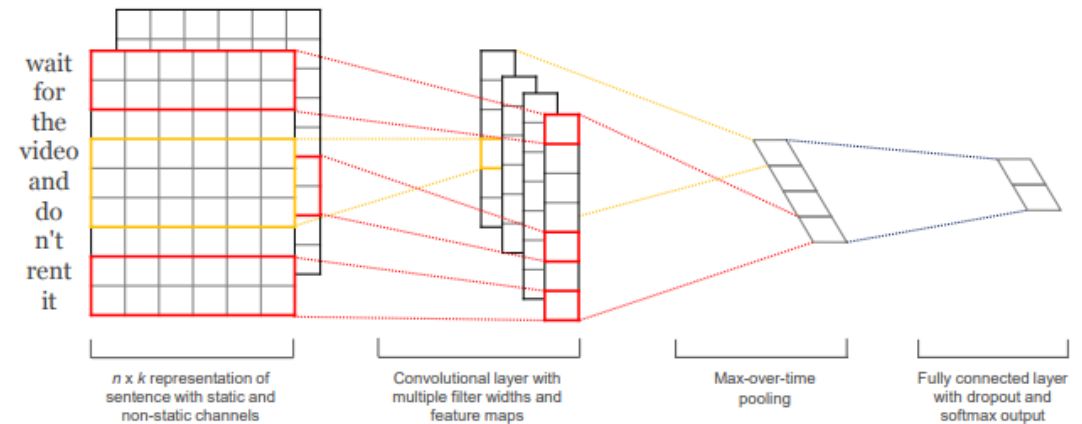
- Convolution operation



- Convolutions over full embeddings!
- Filter size  $n \times m$  ( $n$  = variable,  $m$  = length of the embeddings)

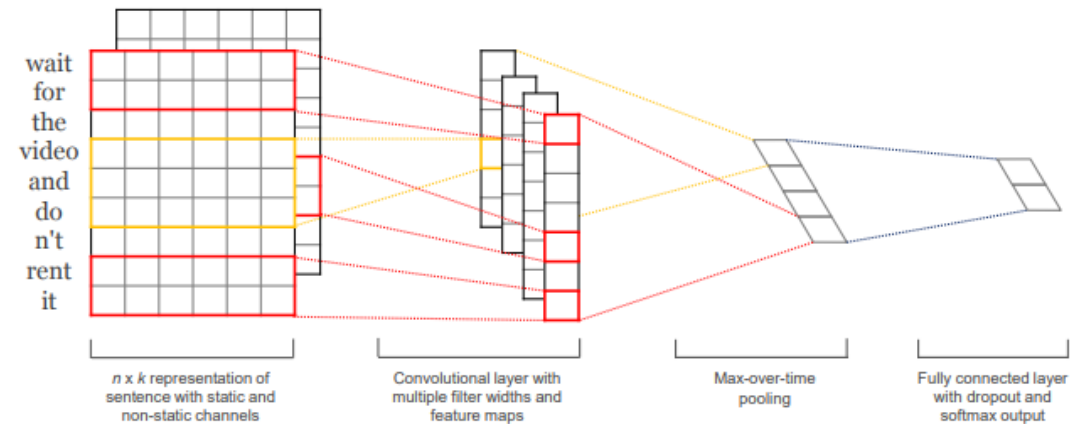
# CNNs for Sentence Classification

- Network architecture
  - No stacked convolutional layers!
  - One layer of convolutions
  - Different filter sizes:  
 $n \in \{3,4,5\}$
  - Each with 100 filters
  - Stride 1 -> Do not skip any words/n-grams

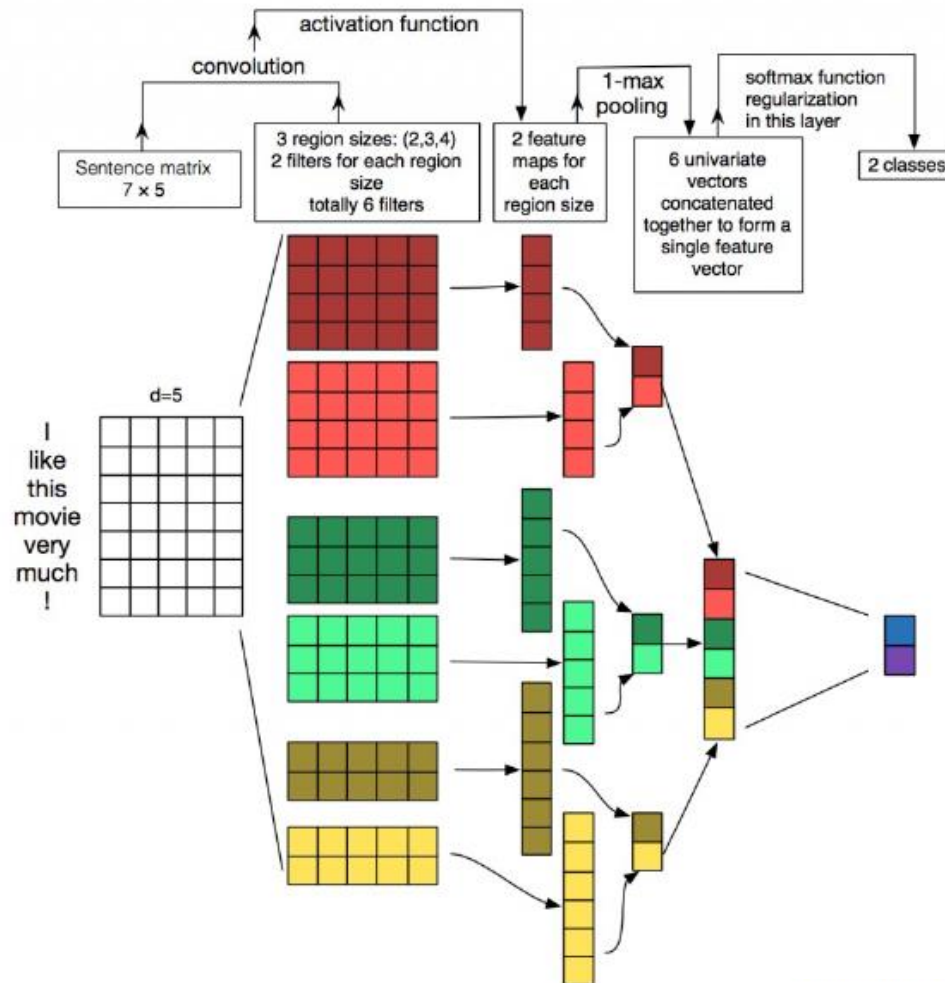


# CNNs for Sentence Classification

- Network architecture
  - Max-over-time pooling
  - Max Pooling over the full filter output -> Select the highest output for each filter
  - Fully connected layer
  - Dropout
  - Softmax



# CNNs for Sentence Classification



<https://dennybritz.com/posts/wildml/understanding-convolutional-neural-networks-for-nlp/>

# CNNs for Sentence Classification: Embeddings

---

- Word embeddings used to encode the input
- Multiple variants:
  - Randomly initialize embeddings, train with model (cnn-rand)
  - Initialize word embeddings with Word2Vec, keep fixed (cnn-static)
  - Initialize word embeddings with Word2Vec, train with model (cnn-nonstatic)
- Finding: Using pre-trained embeddings and further optimizing them for the task at hand often works best!



# CNNs for Sentence Classification: Embeddings

- Trick: Multi-channel word embeddings (cnn-multichannel)
  - Represent input as 3-dimensional matrix
  - 3rd dimension: different, independent word embedding vectors
  - Convolutions computed slightly differently, general ideas still apply
- During training:
  - keep one dimension fixed
  - Train the other
- Effect: Keep information from original embeddings, but also include „extra“ for the task/dataset

	$w_{0,0,0}$	$w_{0,1,0}$	...	$w_{0,m,0}$
$v$	$w_{0,0,1}$	$w_{0,1,1}$	...	$w_{0,m,1}$
	$w_{1,0,1}$	$w_{1,1,1}$	...	$w_{1,m,1}$
$v$				
	$w_{n,0,1}$	$w_{n,1,1}$	...	$w_{n,m,1}$

# CNNs for Sentence Classification: Regularization

---

- Multiple regularization methods used:
  - Dropout
    - Apply Dropout after the fully connected layer
  - L2-maxnorm constraint
    - Set a hard limit  $s$  on l2-norm for weights  $w$  of the fully connected layer
    - If, after the update step,  $\|w\|_2 > s$ , rescale  $w$  to  $\|w\|_2 = s$
    - Prevents single weights from getting too large

# CNNs for Sentence Classification: Evaluation

---

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-rand	76.1	45.0	82.7	89.6	91.2	79.8	83.4
CNN-static	81.0	45.5	86.8	93.0	92.8	84.7	<b>89.6</b>
CNN-non-static	<b>81.5</b>	48.0	87.2	93.4	93.6	84.3	89.5
CNN-multichannel	81.1	47.4	<b>88.1</b>	93.2	92.2	<b>85.0</b>	89.4
RAE (Socher et al., 2011)	77.7	43.2	82.4	—	—	—	86.4
MV-RNN (Socher et al., 2012)	79.0	44.4	82.9	—	—	—	—
RNTN (Socher et al., 2013)	—	45.7	85.4	—	—	—	—
DCNN (Kalchbrenner et al., 2014)	—	48.5	86.8	—	93.0	—	—
Paragraph-Vec (Le and Mikolov, 2014)	—	<b>48.7</b>	87.8	—	—	—	—
CCAE (Hermann and Blunsom, 2013)	77.8	—	—	—	—	—	87.2
Sent-Parser (Dong et al., 2014)	79.5	—	—	—	—	—	86.3
NBSVM (Wang and Manning, 2012)	79.4	—	—	93.2	—	81.8	86.3
MNB (Wang and Manning, 2012)	79.0	—	—	<b>93.6</b>	—	80.0	86.3
G-Dropout (Wang and Manning, 2013)	79.0	—	—	93.4	—	82.1	86.1
F-Dropout (Wang and Manning, 2013)	79.1	—	—	<b>93.6</b>	—	81.9	86.3
Tree-CRF (Nakagawa et al., 2010)	77.3	—	—	—	—	81.4	86.1
CRF-PR (Yang and Cardie, 2014)	—	—	—	—	—	82.7	—
SVM <sub>S</sub> (Silva et al., 2011)	—	—	—	—	<b>95.0</b>	—	—

# CNNs for Sentence Classification: Evaluation

---

- General findings:
  - Model performs very well!
  - Using pre-trained word embeddings is better than randomly initializing
  - No clear advantage for using cnn-static, cnn-nonstatic or cnn-multichannel

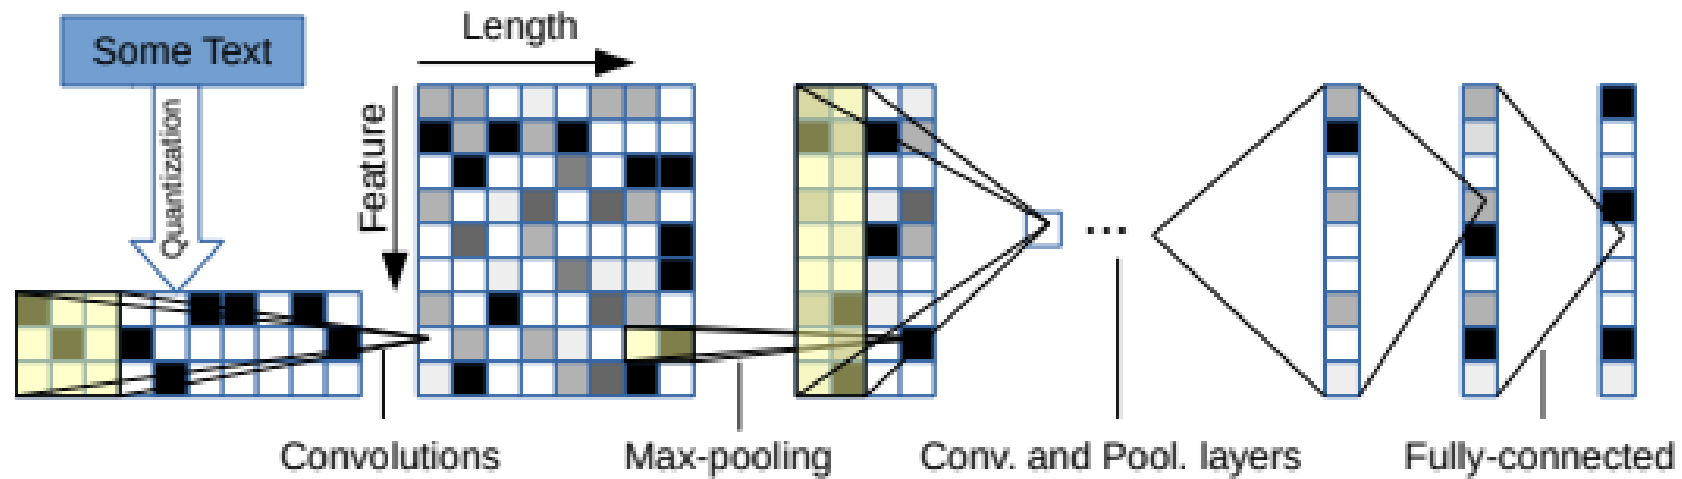
# Character level CNNs

---

- Idea: Learn from scratch, ignore notion of words (!)
- Instead of having a sequence of words for which we train embeddings: Train on sequence of characters with embeddings
- Alphabet used (70 different characters):

```
abcdefghijklmnopqrstuvwxyz0123456789,;.:'''/\|_@#$$%^&*~+-=<>()[]{} 
```

# Character level CNNs




# Training

---

- The algorithm used is stochastic gradient descent (SGD) with a minibatch of size 128, using momentum 0.9 and initial step size 0.01 which is halved every 3 epochs for 10 times.
- Problem: Text length
- Solution: Text a fixed length of 1014 characters (!). If necessary, padding with spaces

# Evaluation

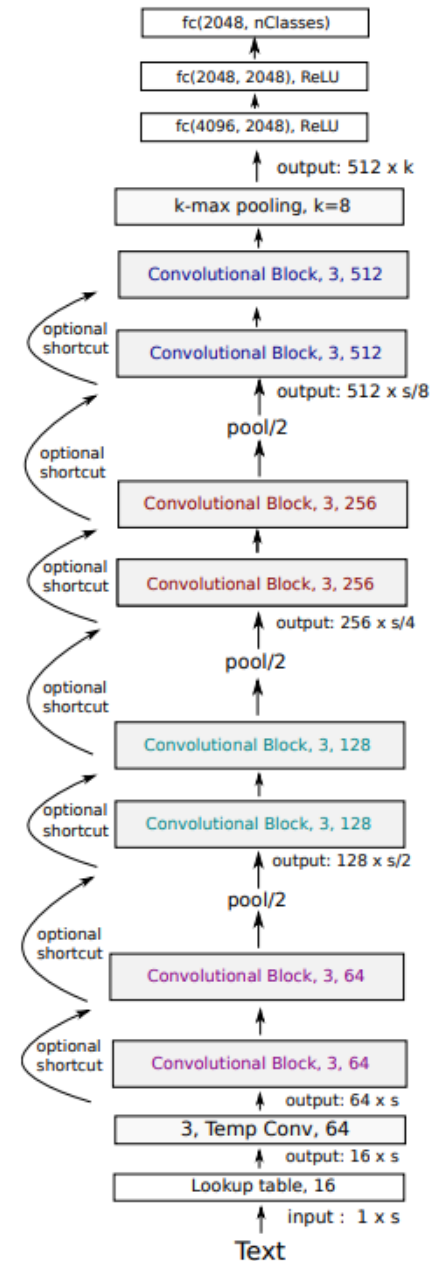
Model	AG	Sogou	DBP.	Yelp P.	Yelp F.	Yah. A.	Amz. F.	Amz. P.
BoW	11.19	7.15	3.39	7.76	42.01	31.11	45.36	9.60
BoW TFIDF	10.36	6.55	2.63	6.34	40.14	28.96	44.74	9.00
ngrams	7.96	2.92	1.37	4.36	43.74	31.53	45.73	7.98
ngrams TFIDF	7.64	2.81	1.31	4.56	45.20	31.49	47.56	8.46
Bag-of-means	16.91	10.79	9.55	12.67	47.46	39.45	55.87	18.39
LSTM	13.94	4.82	1.45	5.26	41.83	29.16	40.57	6.10
Lg. w2v Conv.	9.92	4.39	1.42	4.60	40.16	31.97	44.40	5.88
Sm. w2v Conv.	11.35	4.54	1.71	5.56	42.13	31.50	42.59	6.00
Lg. w2v Conv. Th.	9.91	-	1.37	4.63	39.58	31.23	43.75	5.80
Sm. w2v Conv. Th.	10.88	-	1.53	5.36	41.09	29.86	42.50	5.63
Lg. Lk. Conv.	8.55	4.95	1.72	4.89	40.52	29.06	45.95	5.84
Sm. Lk. Conv.	10.87	4.93	1.85	5.54	41.41	30.02	43.66	5.85
Lg. Lk. Conv. Th.	8.93	-	1.58	5.03	40.52	28.84	42.39	5.52
Sm. Lk. Conv. Th.	9.12	-	1.77	5.37	41.17	28.92	43.19	5.51
Lg. Full Conv.	9.85	8.80	1.66	5.25	38.40	29.90	40.89	5.78
Sm. Full Conv.	11.59	8.95	1.89	5.67	38.82	30.01	40.88	5.78
Lg. Full Conv. Th.	9.51	-	1.55	4.88	38.04	29.58	40.54	5.51
Sm. Full Conv. Th.	10.89	-	1.69	5.42	37.95	29.90	40.53	5.66
Lg. Conv.	12.82	4.88	1.73	5.89	39.62	29.55	41.31	5.51
Sm. Conv.	15.65	8.65	1.98	6.53	40.84	29.84	40.53	5.50
Lg. Conv. Th.	13.39	-	1.60	5.82	39.30	28.80	40.45	4.93
Sm. Conv. Th.	14.80	-	1.85	6.49	40.16	29.84	40.43	5.67


Small
Large



# Even deeper CNNs

- Same ideas, deeper architecture
- Upto 49 layers of blocks
- Achieves better performance
- High computational cost



# Reference and further reading

---

- <https://cs231n.github.io/convolutional-networks/>
- <https://www.youtube.com/watch?v=bNb2fEVKeEo>