

Assignment: Convolutional Neural Networks for Text Mining

Course: Text Mining (Summer Term 2025)

Instructor: [Your Name]

Submission Deadline: [Insert Date]

Part A: Conceptual Questions (Short Answer)

1. Explain the difference between **word embeddings** and **character-level representations**. In which cases might one outperform the other?
 2. Why are CNNs more parameter-efficient than fully connected networks? Illustrate with an example.
 3. Describe the role of **stride** and **padding** in convolution operations. How do they affect the output feature map dimensions?
 4. Discuss the intuition behind **max pooling** in CNNs. Why is it especially useful for text classification tasks?
 5. Compare and contrast the following embedding strategies:
 - cnn-rand
 - cnn-static
 - cnn-nonstatic
 - cnn-multichannel
-

Part B: Applied Exercises (Programming)

Use **Python** with **PyTorch** or **TensorFlow**. Submit your code along with explanations.

1. Word-level CNN for Sentiment Classification

- Train a CNN-based classifier on the IMDB sentiment dataset.
- Experiment with at least two embedding setups:
 - Pre-trained Word2Vec or GloVe embeddings (fixed).
 - Trainable embeddings (initialized randomly).
- Compare results and discuss your findings.

2. Effect of Convolution Parameters

- Implement a toy CNN where you vary **filter size** (e.g., 2, 3, 5 words) and **stride** values.
- Visualize how these changes affect the extracted n-gram features.

3. Character-level CNN

- Implement a simple character-level CNN for text classification (you may use a subset of data if training is heavy).
- Reflect on performance compared to the word-level CNN.

Part C: Critical Reflection (Essay, 800–1000 words)

- Evaluate the impact of CNNs in **NLP** compared to earlier models (bag-of-words, n-grams, RNNs).
- Discuss the limitations of CNNs for text classification. Where might RNNs or Transformers still be preferable?
- Suggest one novel research direction or modification to CNNs that could improve their performance on modern NLP tasks.